

**COMP0087 25/26**

# **Lecture 1: Introduction**

**13/01/2026**

# **Who are we?**

# Instructors

Yao Lu

Lecturer in Natural Language Processing

Contact: [yao.lu@cs.ucl.ac.uk](mailto:yao.lu@cs.ucl.ac.uk)

Research: language model pretraining, prompt optimisation

# Instructors

Ehsan Shareghi

Associate Professor in Natural Language Processing

Contact: [ehsan.shareghi@gmail.com](mailto:ehsan.shareghi@gmail.com)

Research: LLM safety, speech processing

# Instructors

Pontus Stenetorp

Professor in Natural Language Processing

Contact: [p.stenetorp@cs.ucl.ac.uk](mailto:p.stenetorp@cs.ucl.ac.uk)

Research: Anything that fits into NLP (Multilinguality, data, knowledge, generalisation, etc)

# Instructors

Vasileios (Bill) Lampos

Associate Professor

Contact: [v.lampos@ucl.ac.uk](mailto:v.lampos@ucl.ac.uk)

Research: Sequence modelling (time series forecasting), LLMs training dynamics (e.g., mode collapse)

# Assumptions about you

You already know the basics of applied ML:

- What the training/dev/test cycle is
- What a loss function is
- What an optimization technique does (at a high level)
- How to evaluate a model (given a standard benchmark)

# Assumptions about you

You can already do these things:

- Perform basic data manipulation
- Build basic knowledge in Python

# Who are you?

Level of experience in NLP

1. Beginner (no prior NLP experience)
2. Intermediate (coursework or projects)
3. Advanced (research or professional experience)

# Who are you?

I'm taking this course to

1. Prepare for job searching
2. Learn how to conduct NLP research
3. Apply NLP to ongoing projects and research
4. Other

# Module communication

# General, non-personal, queries

\*Largest\* postgraduate module

General, non-personal, queries:

Please, pretty please use the Moodle forum

# Personal queries

Personal queries:

[comp0087@cs.ucl.ac.uk](mailto:comp0087@cs.ucl.ac.uk)

# General project-related queries

General, project-related queries:

Moodle forum

# Your project-related queries

Specific, project-related queries:

Assigned teaching assistant

# Administrative-related queries

Administrative logistics (registration, etc.):

[cs.pgt-students@ucl.ac.uk](mailto:cs.pgt-students@ucl.ac.uk)

# Module content

# Learning goals

1. Foundational “building blocks” for LLMs
2. Advanced LLM development/applications
3. Practical, empirical understanding of building NLP systems
  - Analysis, empirical experiments, etc.
  - Group project

# Teaching plan

Basic NLPs 3-4 weeks

Advanced NLPs (LLMs and applications) 3-4 weeks

Topic Lectures from Industry speakers (e.g., Meta AI, Microsoft etc) 2-3 weeks

NLP Course Projects (supervised by our awesome TAs and faculty members) 1-10 weeks

# Teaching plan: Foundations

Tokenisation

Token  
representation

What is your input?

In what form?

# Teaching plan: Foundations

Traditional LM

Neural LM

How to model the data?

# Teaching plan: Foundations

Attention

Transformer

Which architecture and why?

# Teaching plan: Foundations

## Prompting

How to "talk" to the model?

# Teaching plan: Advanced Topics

## Advanced NLPs

- Tuning Large Language Models
- Alignment with Human Feedback
- Reasoning LLMs
- Retrieval Augmented Generation (RAG)
- LLMs Pretraining

...

# Teaching plan: Industry Lectures

Please let us know which topics you're interested in.

# Group project

# Assessment

Coursework (Group): 100% of module mark

- Mandatory bi-weekly progress meetings
- Single-page progress reports (due date TBD)
- Eight-page final report
- Not including references

# Group formation

- Groups of \*five\* or \*six\*
- \* Exceptions will be rare
- Group formation due 18 January by midnight
- \* Start looking \*now\*
- \* Decide on group name
- \* Post UCL e-mails in dedicated Moodle forum

# Project requirements

- Must involve language
- Empirical investigation
- Designed by students, with the \*help\* of us
- \* Ultimate responsibility and control is with \*you\*
- Outcomes \*may\* be publishable
- \* But \*no\* guarantees

# Project timeline sketch

- Week 1 to 2: Project design and literature review
- Week 3 to 4: Data collection, experimental design, initial implementation
- Week 5 to 6: Experiments not working out: Revisions, revisions, revisions,
- Week 7 to 8: Additional experiments and analysis
- Week 9 to 10: Report writing

# Project design

- Minimum viable project
- Align with group interests and teaching assistant expertise
- Design to be "split"

# Computational resources

- \*Limited\*: Design your projects with this in mind!
- - UCL Computer Science:
- \* <https://tsg.cs.ucl.ac.uk/gpus>
- \* ~30 x NVIDIA RTX 4060 cards (16GB)
- \* ~25 x NVIDIA RTX 3090 cards (24GB)

# Computational resources

- Google Colab:
- \* <https://colab.research.google.com>
- \* Web "notebooks" with free GPUs and TPUs (LLM access?)
- \* Can pay for better service (GPU server rental)
- co:here credits
- \* Application process announced shortly

# Project design example

I want to train a language model.

I want to train a language model that better  
understands twitter posts.

# Project design example

- What's broken? (Problem)
- How will you know it's fixed? (Metrics)
- What will you try? (Approach)
- Can you actually do it? (Feasibility)
- What if it fails? (Contingency)

# Project design example

## What's Broken?

- Why do current LLMs struggle with tweets?
- What evidence supports this? (Literature review)
- Show examples of failures

# Project design example

## How Will You Know It's Fixed?

- Which task? (sentiment, NER, sarcasm detection, etc.)
- Which dataset? (TweetEval? SemEval? Custom?)
- What's the baseline? (GPT-4? RoBERTa-Twitter? Prompting vs. fine-tuning?)

# Project design example

## What Will You Try?

- Which base model?
- Prompting, fine-tuning, or pretraining? (Keep budget in mind)
- What training data? (Source, size, licensing)
- Why should this approach work? (Hypothesis)

# Project design example

## Can You Actually Do It?

- GPU hours (check maximum available resources)
- Training data (size, cost)
- Evaluation data (quality, cost)
- Timeline (10 weeks in total)
- Split tasks

# Project design example

## What If It Fails?

- No improvement: Why? Data quality? Wrong task?  
Write negative result analysis
- Reframe as benchmark or dataset contribution
- Discuss with your TA in advance

# Teaching assistants

Adam Oomerjee-Vawda ([adam.vawda@gmail.com](mailto:adam.vawda@gmail.com))

- PhD student at UCL AI centre
- Interests: Efficient models, reasoning
- Favourite paper:
- Bottlenecked Transformers: Periodic KV Cache Abstraction for Generalised Reasoning. <https://arxiv.org/abs/2505.16950>

Edan Toledo (edan.toledo.24@ucl.ac.uk)

- PhD student at UCL and Facebook
- Interests:
  - \* General search methods \* Agentic systems \* Meta-learning
  - \* Reinforcement learning
- Favourite paper:
  - \* "A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play" by Silver et al. (2018)
  - \* <https://www.science.org/doi/10.1126/science.aarr6404>

Eduardo Sánchez (eduardo.sanchez.22@ucl.ac.uk)

- final-year PhD student at UCL and Facebook
- Interests:
- \* Low-resource languages \* Multilinguality
- Favourite paper:
- \* "Linguini: A benchmark for language-agnostic linguistic reasoning" by Sánchez et al. (2024) <https://arxiv.org/abs/2409.12126>

Jiayi Wang (jiayi.lin.wang@ucl.ac.uk)

- final-year PhD student
- Interests:
  - \* Inclusive and efficient multilingual large language models
- Favourite paper:
  - \* "Multilingual Pretraining Using a Large Corpus Machine-Translated from a Single Source Language" by Wang et al. (2023)
    - \* <https://arxiv.org/abs/2410.23956>

Karen Hambardzumyan (karen.hambardzumyan.23@ucl.ac.uk)

- PhD student at UCL and Facebook
- Interests:
  - \* Mechanistic interpretability \* Multi-agent systems
- Favourite paper:
  - \* "WARP: Word-level Adversarial ReProgramming" by Hambardzumyan et al. (2021) \* <https://aclanthology.org/2021.acl-long.381>

Lovish Madaan (lovish.madaan.23@ucl.ac.uk)

- PhD student at UCL and Facebook (GenAI)
- Interests:
  - \* Generalisation behaviours in reinforcement learning \*
  - Better/harder evaluation
- Favourite paper:
  - \* "Amortizing intractable inference in large language models" by Hu et al. (2024) \* <https://arxiv.org/abs/2310.04363>

Ralph Tang (r-tang.25@ucl.ac.uk)

- PhD student at NLP group
- Interests: \* Multimodality \* LLM/VLM interpretability
- Favourite paper:
- What the DAAM: Interpreting Stable Diffusion Using Cross Attention. ACL 2023

Hossein A (Saeed) Rahmani  
(hossein.rahamani.22@ucl.ac.uk)

- PhD student at NLP group
- Interests: \* Evaluations and benchmarks \* Self-Improvement
- Favourite paper:
- "Large Language Models Cannot Self-Correct Reasoning Yet" by Huang et al. (2024) <https://arxiv.org/abs/2310.01798>

# Tokenisation

Let's study NLP together.



Human (Word-level)

Let's

study

NLP

together

.

Total: 5 tokens (words and punctuation)



GPT3.5 (Subword-level Tokenization)

Let

' s

study

N

LP

together

.

Total: 7 tokens (subword units for better vocabulary coverage)

# From Text to Tokens

Tokens	Characters
7	25

Let's study NLP together.

[10267, 596, 4007, 452, 12852, 3871, 13]

Text

Token IDs

Text

Token IDs

# The Reverse Spelling Problem

strawberry -> y-r-r-e-b-w-a-r-t-s

Why is this difficult?

# The Reverse Spelling Problem

strawberry -> [“str”, “aw”, “berry”]

-> y-r-r-e-b-w-a-r-t-s?

Why is this difficult?

# The Reverse Spelling Problem

strawberry -> [“str”, “aw”, “berry”]  
-> [“yr”, “reb”, “warts”]

Why is this difficult?

# The Reverse Spelling Problem

strawberry -> [496, 675, 15717]  
-> [11160, 32575, 64156]

Why is this difficult?

# The Number Comparison Problem

"Is 9.8 greater than 9.11?"

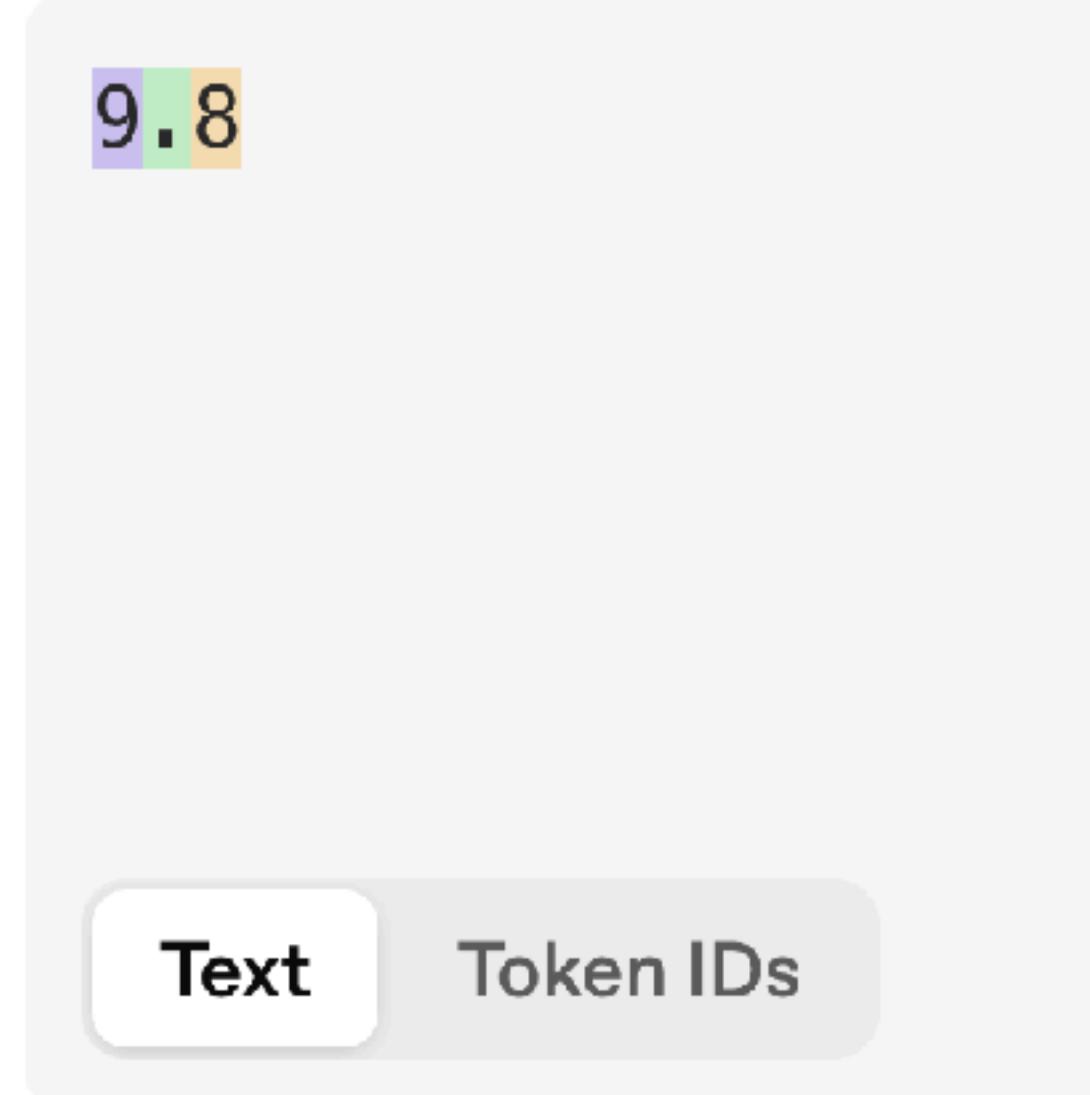
# The Number Comparison Problem

Human: 9.**8** > 9.**1**1

# The Number Comparison Problem

GPT4:      9.8      <      9.11

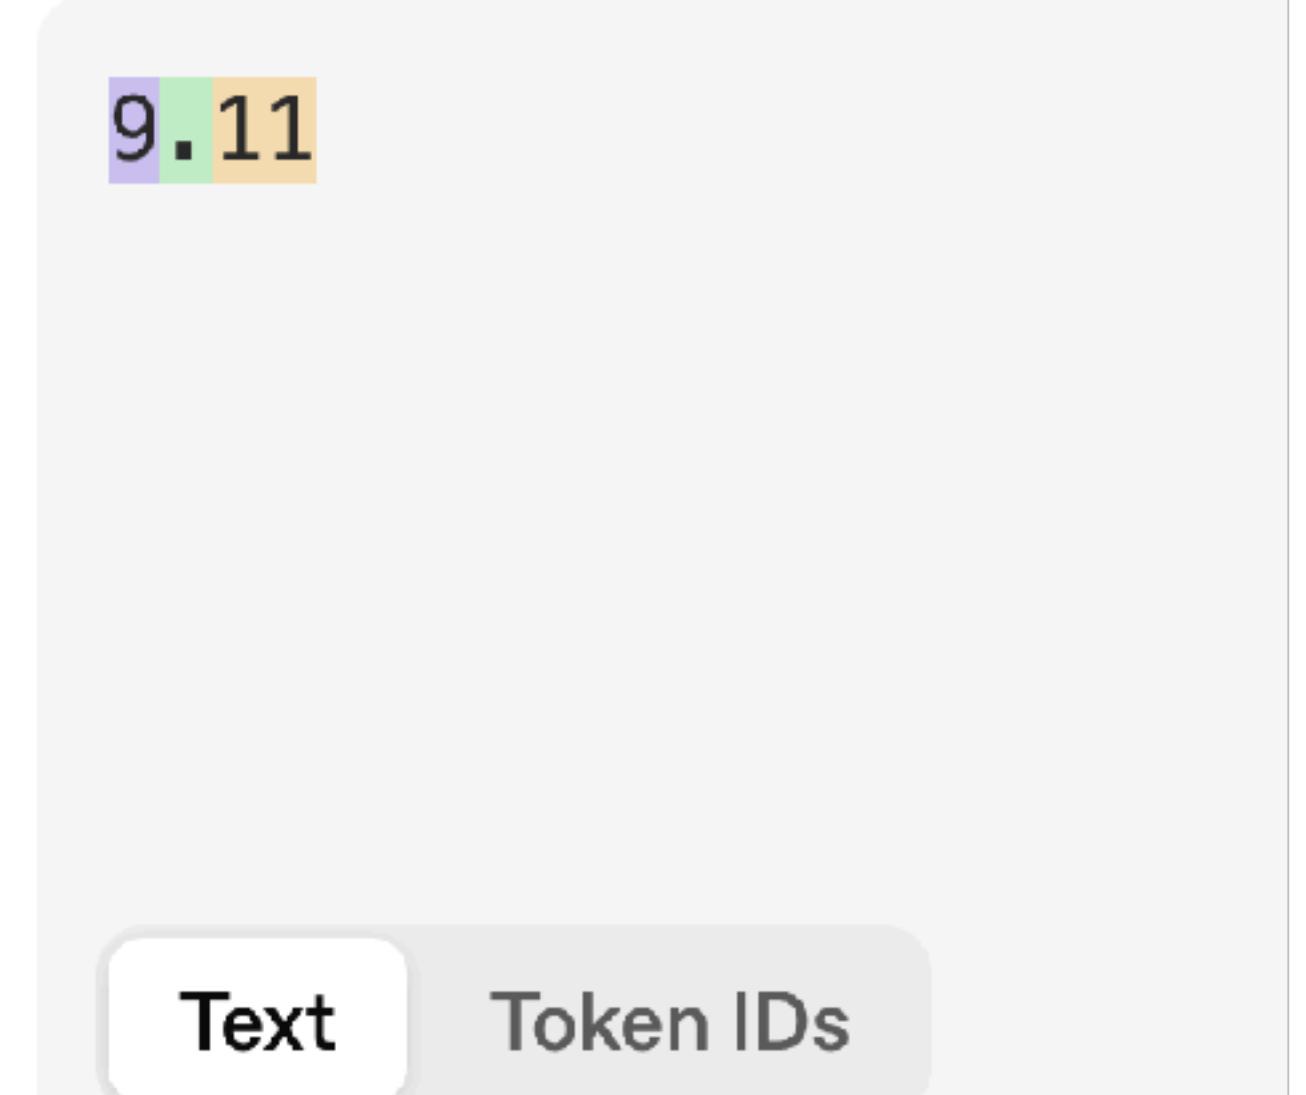
Tokens	Characters
3	3



9.8

Text      Token IDs

Tokens	Characters
3	4



9.11

Text      Token IDs

# Definition of tokenisation

- The process of splitting **text** into **meaningful tokens**, which are then mapped to **numerical IDs**.

# Classical tokenisation methods

Character-level tokenisation: strawberry -> [s, t, r, a, w, b, e, r, r, y]

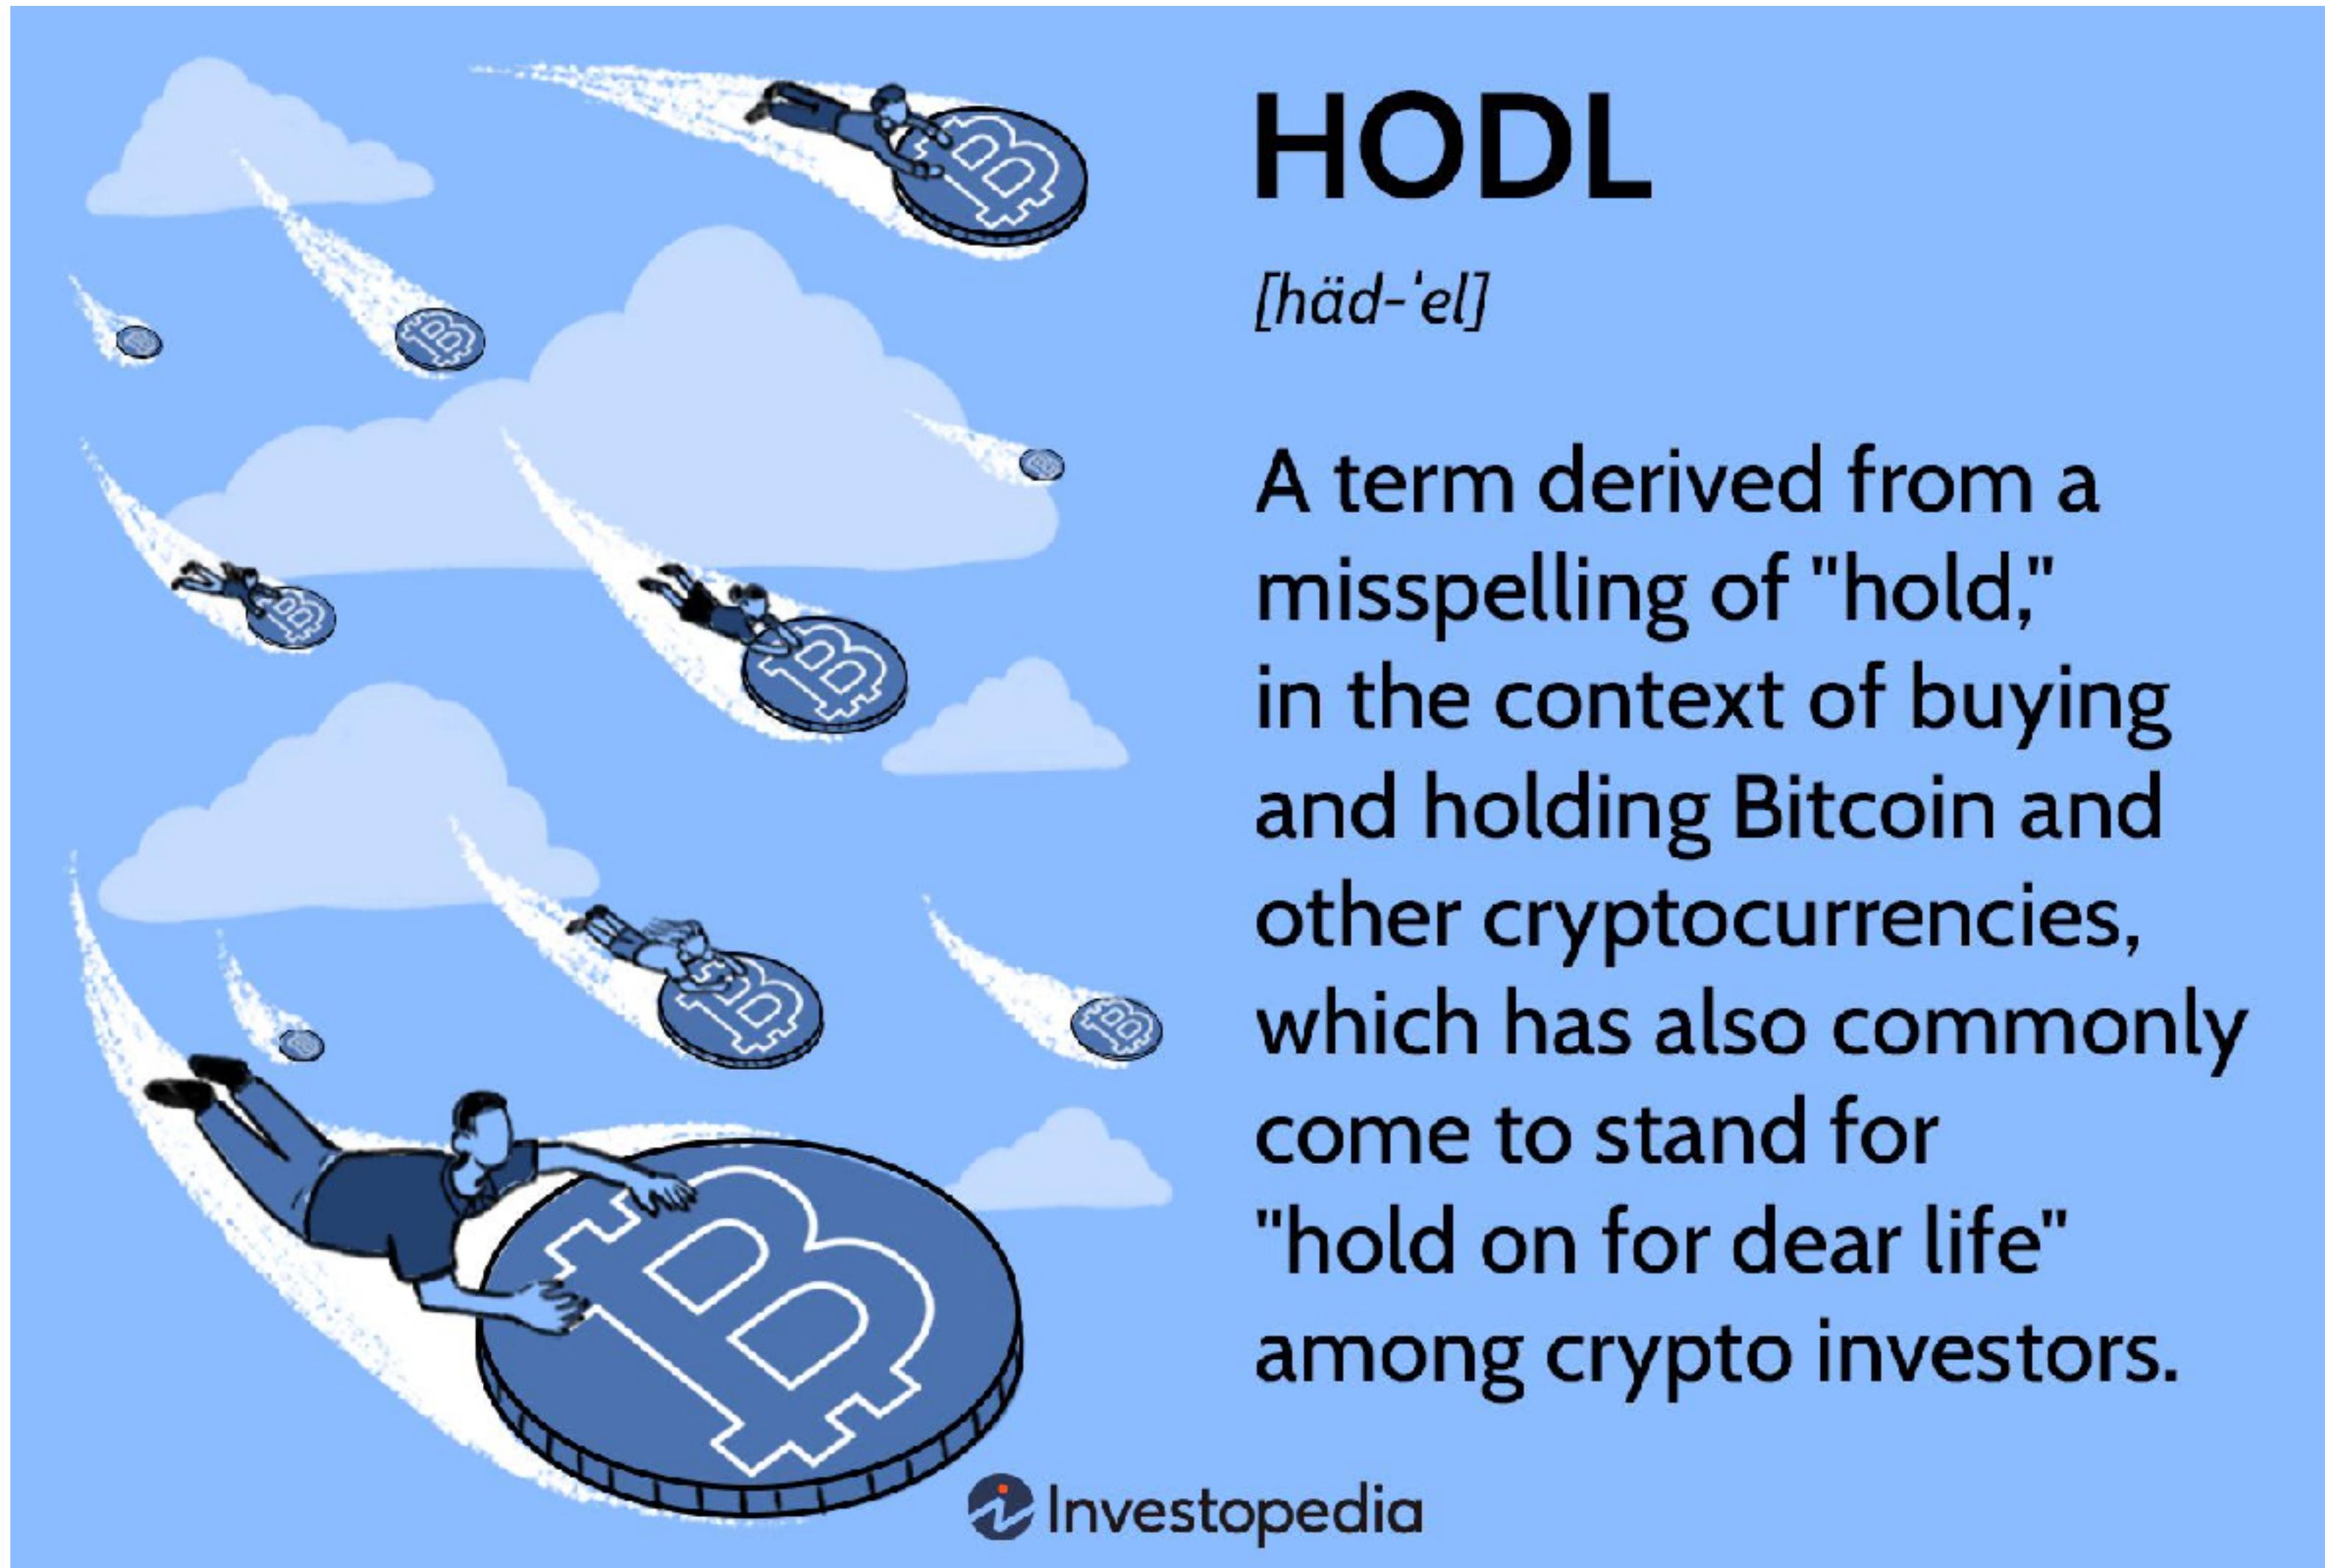
- Pros: a very small set of symbols (100?) can represent almost everything
- Cons: less meaningful segmentation

# Classical tokenisation methods

Word-level tokenisation: “I am happy” -> [Index(I), Index(am), Index(happy)]

- Pros: meaningful and natural segmentation (by whitespace)
- Cons: huge vocabulary space

# Out of Vocabulary Tokens



# Out of Vocabulary Tokens

Solution: replace out of vocabulary tokens with <UNK>

“My favourite food is <UNK>”

“My favourite <UNK> is <UNK>”

“<UNK> favourite <UNK> is <UNK>”

# Tokenisation design

	Character-based method	Word-level method	?
Vocabulary size	~100	100,000	1k-100k
Corpus coverage	High (any text)	Medium (OOV issues)	High
Meaningfulness	Low	High	High
Sequence length	Very long	Short	Medium

# Subword Tokenisation

- Frequent words stay as single tokens ("cat").
- Rare/complex words split into meaningful sub-units ("tokenisation" -> "token" + "isation").
- Handles OOV: Can decompose new words ("ChatGPT" -> "Chat" + "G" + "PT").

# Byte-Pair Encoding (BPE)

strawberry -> [“str”, “aw”, “berry”]

# Byte-Pair Encoding (BPE)

“London”

“Londres”

“Londra”

# Byte-Pair Encoding (BPE)

“London”

“Londres”

“Londra”

{L, o, n, d, r, e, s, a}

# Byte-Pair Encoding (BPE)

“London”

(L, o) 3

“Londres”

“Londra”

{L, o, n, d, r, e, s, a}

# Byte-Pair Encoding (BPE)

“London” (L, o) 3

“Londres” (o, n) 4

“Londra”

{L, o, n, d, r, e, s, a}

# Byte-Pair Encoding (BPE)

“London” (L, o) 3

“Londres” (o, n) 4  
(n, d) 3

“Londra”

{L, o, n, d, r, e, s, a}

# Byte-Pair Encoding (BPE)

“London”	(L, o) 3
“Londres”	(o, n) 4
“Londra”	(n, d) 3 (d, o) 1

{L, o, n, d, r, e, s, a}

# Byte-Pair Encoding (BPE)

“London”	(L, o) 3
“Londres”	(o, n) 4
“Londra”	(n, d) 3
{L, o, n, d, r, e, s, a}	(d, o) 1
	(d, r) 2

# Byte-Pair Encoding (BPE)

“London”	(L, o) 3
“Londres”	(o, n) 4
“Londra”	(n, d) 3
{L, o, n, d, r, e, s, a}	(d, o) 1
	(d, r) 2
	(r, e) 1

# Byte-Pair Encoding (BPE)

“London”	(L, o) 3
“Londres”	(o, n) 4
“Londra”	(n, d) 3
{L, o, n, d, r, e, s, a}	(d, o) 1
	(d, r) 2
	(r, e) 1
	(e, s) 1

# Byte-Pair Encoding (BPE)

“London”	(L, o) 3
“Londres”	(o, n) 4
“Londra”	(n, d) 3
{L, o, n, d, r, e, s, a}	(d, o) 1
	(d, r) 2
	(r, e) 1
	(e, s) 1
	(r, a) 1

# Byte-Pair Encoding (BPE)

“London”	(L, o) 3
“Londres”	(o, n) 4
“Londra”	(n, d) 3
{L, o, n, d, r, e, s, a, on}	(d, o) 1
rule #1: (o, n) -> on	(d, r) 2
	(r, e) 1
	(e, s) 1
	(r, a) 1

# Byte-Pair Encoding (BPE)

“L on d on”

“L on d r e s”

“L on d r a”

{L, o, n, d, r, e, s, a, on}

rule #1: (o, n) -> on

# Byte-Pair Encoding (BPE)

“L on d on”

(L, on) 3

“L on d r e s”

“L on d r a”

{L, o, n, d, r, e, s, a, on}

rule #1: (o, n) -> on

# Byte-Pair Encoding (BPE)

“L on d on”

(L, on) 3

“L on d r e s”

(on, d) 3

“L on d r a”

{L, o, n, d, r, e, s, a, on}

rule #1: (o, n) -> on

# Byte-Pair Encoding (BPE)

“L on d on” (L, on) 3

“L on d r e s” (on, d) 3

“L on d r a” (d, on) 1

{L, o, n, d, r, e, s, a, on}

rule #1: (o, n) -> on

# Byte-Pair Encoding (BPE)

“L on d on” (L, on) 3

“L on d r e s” (on, d) 3

“L on d r a” (d, on) 1

(d, r) 2

{L, o, n, d, r, e, s, a, on}

rule #1: (o, n) -> on

# Byte-Pair Encoding (BPE)

“L on d on” (L, on) 3

“L on d r e s” (on, d) 3

“L on d r a” (d, on) 1

(d, r) 2

{L, o, n, d, r, e, s, a, on} (r, e) 1

rule #1: (o, n) -> on

# Byte-Pair Encoding (BPE)

“L <u>on</u> d <u>on</u> ”	(L, on) 3
“L <u>on</u> d r e s”	(on, d) 3
“L <u>on</u> d r a”	(d, on) 1
	(d, r) 2
{L, o, n, d, r, e, s, a, on}	(r, e) 1
rule #1: (o, n) -> on	(e, s) 1

# Byte-Pair Encoding (BPE)

“L on d on” (L, on) 3

“L on d r e s” (on, d) 3

“L on d r a” (d, on) 1

(d, r) 2

{L, o, n, d, r, e, s, a, on} (r, e) 1

rule #1: (o, n) -> on (e, s) 1

(r, a) 1

# Byte-Pair Encoding (BPE)

“L on d on” (L, on) 3

“L on d r e s” (on, d) 3

“L on d r a” (d, on) 1

(d, r) 2

{L, o, n, d, r, e, s, a, on, **Lon**} (r, e) 1

(e, s) 1

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon (r, a) 1

# Byte-Pair Encoding (BPE)

“Lon d on”

“Lon d r e s”

“Lon d r a”

{L, o, n, d, r, e, s, a, on, Lon}

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

# Byte-Pair Encoding (BPE)

“Lon d on”

(Lon, d) 3

“Lon d r e s”

“Lon d r a”

{L, o, n, d, r, e, s, a, on, Lon}

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

# Byte-Pair Encoding (BPE)

“Lon d on”

(Lon, d) 3

“Lon d r e s”

(d, on) 1

“Lon d r a”

{L, o, n, d, r, e, s, a, on, Lon}

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

# Byte-Pair Encoding (BPE)

“Lon d on”

(Lon, d) 3

“Lon d r e s”

(d, on) 1

“Lon d r a”

(d, r) 2

{L, o, n, d, r, e, s, a, on, Lon}

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

# Byte-Pair Encoding (BPE)

“Lon d on”

(Lon, d) 3

“Lon d r e s ”

(d, on) 1

“Lon d r a”

(d, r) 2

{L, o, n, d, r, e, s, a, on, Lon}

(r, e) 1

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

# Byte-Pair Encoding (BPE)

“Lon d on”

(Lon, d) 3

“Lon d r e s”

(d, on) 1

“Lon d r a”

(d, r) 2

{L, o, n, d, r, e, s, a, on, Lon}

(r, e) 1

(e, s) 1

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

# Byte-Pair Encoding (BPE)

“Lon d on”

(Lon, d) 3

“Lon d r e s ”

(d, on) 1

“Lon d r a”

(d, r) 2

{L, o, n, d, r, e, s, a, on, Lon}

(r, e) 1

(e, s) 1

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

(r, a) 1

# Byte-Pair Encoding (BPE)

“Lon d on”

(Lon, d) 3

“Lon d r e s”

(d, on) 1

“Lon d r a”

(d, r) 2

{L, o, n, d, r, e, s, a, on, Lon, **Lond**}

(r, e) 1

rule #1: (o, n) -> on

(e, s) 1

rule #2: (L, on) -> Lon

(r, a) 1

rule #3: (Lon, d) -> Lond

# Byte-Pair Encoding (BPE)

“Lond on”

“Lond r e s ”

“Lond r a”

{L, o, n, d, r, e, s, a, on, Lon, Lond}

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

# Byte-Pair Encoding (BPE)

“Lond on”

(Lond, on) 1

“Lond r e s ”

“Lond r a”

{L, o, n, d, r, e, s, a, on, Lon, Lond}

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

# Byte-Pair Encoding (BPE)

“Lond on”

(Lond, on) 1

“Lond r e s ”

(Lond, r) 2

“Lond r a”

{L, o, n, d, r, e, s, a, on, Lon, Lond}

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

# Byte-Pair Encoding (BPE)

“Lond on”

(Lond, on) 1

“Lond r e s ”

(Lond, r) 2

“Lond r a”

(r, e) 1

{L, o, n, d, r, e, s, a, on, Lon, Lond}

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

# Byte-Pair Encoding (BPE)

“Lond on”

(Lond, on) 1

“Lond r e s ”

(Lond, r) 2

“Lond r a”

(r, e) 1

{L, o, n, d, r, e, s, a, on, Lon, Lond}

(e, s) 1

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

# Byte-Pair Encoding (BPE)

“Lond on”

(Lond, on) 1

“Lond r e s ”

(Lond, r) 2

“Lond r a”

(r, e) 1

{L, o, n, d, r, e, s, a, on, Lon, Lond}

(e, s) 1

rule #1: (o, n) -> on

(r, a) 1

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

# Byte-Pair Encoding (BPE)

“Lond on”

(Lond, on) 1

“Lond r e s ”

(Lond, r) 2

“Lond r a”

(r, e) 1

{L, o, n, d, r, e, s, a, on, Lon, Lond,  
Londr}

(e, s) 1

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

(r, a) 1

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

# Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,  
Lond, Londr}

Vocabulary

rule #1: (o, n) -> on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

Merge Rules

# Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,  
Lond, Londr}

rule #1: (o, n) -> on

**rondLon**

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

# Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,  
Lond, Londr}

rule #1: (o, n) -> on

r on d L on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

# Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,  
Lond, Londr}

rule #1: (o, n) -> on

r on d L on

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

# Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,  
Lond, Londr}

rule #1: (o, n) -> on

r on d Lon

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

# Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,  
Lond, Londr}

rule #1: (o, n) -> on

r on d Lon

rule #2: (L, on) -> Lon

rule #3: (Lon, d) -> Lond

No merges for #3 and #4

rule #4: (Lond, r) -> Londr

# Byte-Pair Encoding (BPE)

{L, o, n, d, r, e, s, a, on, Lon,  
Lond, Londr}

rule #1: (o, n) -> on

rondLon

rule #2: (L, on) -> Lon

**r on d Lon**

rule #3: (Lon, d) -> Lond

rule #4: (Lond, r) -> Londr

# Byte-Pair Encoding (BPE)

{0: L, 1: o, 2:n, 3: d,  
4: r, 5: e, 6: s,  
7: a, 8: on, 9: Lon,  
10: Lond, 11: Londr}

rondLon

r on d Lon

[4, 8, 3, 9]

# Why is BPE important?

“London”	[23421]	[4222]
“Londres”	[43, 623, 411]	[26432]
“Londra”	[43, 623, 430]	[17681, 520]

GPT-2

Mistral

# BPE is data sensitive

90% general web

5% programming

5% math

st raw berry

80% general web

10% programming

10% math

str aw berry

# BPE is data sensitive

90% general web

5% programming

5% math

+ emoji

build tokenizer

general web

programming

math

(remove all emojis)

pretraining data

# BPE is data sensitive

Model	#Tokens	Tied Emb.	#Confirmed	Examples
GPT-2 Medium (0.4B)	50,257	Yes	49/999	InstoreAndOnline reportprint _externalToEVA
GPT-2 XL (1.5B)	50,257	Yes	67/999	InstoreAndOnline _RandomRedditor embedreportprint
GPT-J 6B	50,400	No*	200/999	_attRot _externalToEVA _SolidGoldMagikarp
Phi-2 (2.7B)	50,295	No*	103/999	DragonMagazine _TheNitrome _SolidGoldMagikarp
Pythia 6.7B	50,277	No	14/993	FFIRMED _taxp _affidav
GPT-NeoX 20B	50,277	No	10/993	FFIRMED _taxp _affidav
OLMo v1.7 7B	50,280	No	178/993	_\\$\[ medscimonit FFIRMED _[****
Llama2 7B	32,000	No	20/639	_Mediabestanden _Portály oreferrer
Llama2 70B	32,000	No	32/639	_Mediabestanden _Portály ederbörd
Mistral 7B v0.3	32,000	No	53/637	\uefc0 ]);\\r \&gt;?[< _febbra _uitgen
Mixtral 8x7B	32,000	No	44/637	\uefc0 _/**\\r \&gt;];\\r
Rakuten 7B	48,000	No	66/957	\uefc0 _/**\\r \&gt; febbra 稲田大学
Qwen1.5 32B	151,646	No	2450/2966	_ForCanBeConvertedToF (stypy \$PostalCodesNL
Qwen1.5 72B Chat	151,646	No	2047/2968	_ForCanBeConverted useRalative _typingsJapgolly
StableLM2 12B	100,288	No	138/1997	_ForCanBeConverted \tTokenNameIdentifier _StreamLazy
Llama3 8B	128,256	No	556/2540	_ForCanBeConverted ӦыңНӦыңN _CLIIIK krvldkf 글상위
Llama3 70B	128,256	No	462/2540	\$PostalCodesNL итися ilmaktadir ーション ;\\r\\r\\n

# **Thank you!**